

Homework Exercises on I/O-Efficient Algorithms

The maximum number of points for all exercises is 30. The grade for this homework set is: (number of scored points)/3.

Exercise Set I/O-Efficient I

IO.I-1 (3 points) Consider a machine with 1 GB of internal memory and a disk with the following properties:

- the average time until the read/write head is positioned correctly to start reading or writing data (seek time plus rotational delay) is 12 ms,
- once the head is positioned correctly, we can read or write at a speed of 60 MB/s.

Furthermore, the standard block size used by the operating system to transfer data between the internal memory and the disk is 4 KB. We use this machine to sort a file containing 10^8 elements, where the size of each element is 500 bytes; thus the size of the file is 50 GB. Our sorting algorithm performs roughly $2 \frac{n}{B} \lceil \log_{M/B} \frac{n}{B} \rceil$ I/Os for a set of n elements, where M is the number of elements that fit into the internal memory and B is the number of elements that fit into a block. Here we assume that 1 KB = 10^3 bytes, 1 MB = 10^6 bytes, and 1 GB = 10^9 bytes.

- Compute the total time in hours spent on I/Os by the algorithm when we work with the standard block size of 4 KB.
- Now suppose we force the algorithm to work with blocks of size 1 MB. What is the time spent on I/Os in this case?
- Same question as (ii) for a block size of 250 MB.

IO.I-2 (1.5+1.5 points) A *stack* is a data structure that supports two operations: we can *push* a new element onto the stack, and we can *pop* the topmost element from the stack. We wish to implement a stack in external memory. To this end we maintain an array $A[0..m-1]$ on disk. The value m , which is the maximum stack size, is kept in internal memory. We also maintain the current number of elements on the stack, s , in internal memory. The *push*- and *pop*-operations can now be implemented as follows:

$Push(A, x)$ 1: if $s = m$ then 2: return “stack overflow” 3: else 4: $A[s] \leftarrow x; s \leftarrow s + 1$ 5: end if	$Pop(A)$ 1: if $s = 0$ then 2: return “stack is empty” 3: else 4: return $A[s-1]; s \leftarrow s - 1$ 5: end if
---	--

Assume A is blocked in the standard manner: $A[0 \dots B-1]$ is the first block, $A[B \dots 2B-1]$ is the second block, and so on.

- Suppose we allow the stack to use only a single block from A in internal memory. Show that there is a sequence of n operations that requires $\Theta(n)$ I/Os.
- Now suppose we allow the stack to keep two blocks from A in internal memory. Prove that any sequence of n *push*- and *pop*-operations requires only $O(n/B)$ I/Os.

IO.I-3 (1 + 1 + 2 points) Suppose we are given two $m \times m$ matrices X and Y , which are stored in row-major order in 2-dimensional arrays $X[0..m-1, 0..m-1]$ and $Y[0..m-1, 0..m-1]$. We wish to compute the product $Z = XY$, which is the $m \times m$ matrix defined by

$$Z[i, j] := \sum_{k=0}^{m-1} X[i, k] \cdot Y[k, j],$$

for all $0 \leq i, j < m$. The matrix Z should also be stored in row-major order. Let $n := m^2$. Assume that n is much larger than M and that $M \geq B^2$. In the questions below, make sure you consider both read- and write-operations.

- (i) Analyze the I/O-complexity of the matrix-multiplication algorithm that simply computes the elements of Z one by one, row by row. You may assume LRU is used as replacement policy.
- (ii) Analyze the I/O-complexity of the same algorithm if X and Z are stored in row-major order while Y is stored in column-major order. You may assume LRU is used as replacement policy.
- (iii) Consider the following alternative algorithm. For a square matrix Q with more than one row and column, let Q_{TL} , Q_{TR} , Q_{BL} , Q_{BR} be the top left, top right, bottom left, and bottom right quadrant, respectively, that result from cutting Q between rows $\lfloor n/2 \rfloor$ and $\lfloor n/2 \rfloor + 1$, and between columns $\lfloor n/2 \rfloor$ and $\lfloor n/2 \rfloor + 1$. We can compute $Z = XY$ recursively by observing that

$$\begin{aligned} Z_{TL} &= X_{TL}Y_{TL} + X_{TR}Y_{BL}, \\ Z_{TR} &= X_{TL}Y_{TR} + X_{TR}Y_{BR}, \\ Z_{BL} &= X_{BL}Y_{TL} + X_{BR}Y_{BL}, \\ Z_{BR} &= X_{BL}Y_{TR} + X_{BR}Y_{BR}. \end{aligned}$$

Analyze the I/O-complexity of this recursive computation. You may assume that m is a power of two and that an optimal replacement policy is used.

Hint: Express the I/O-complexity as a recurrence with a suitable base case, and solve the recurrence.

Exercise Set I/O-Efficient II

IO.II-1 (1+1+1 points) Consider an image-processing application that repeatedly scans an image, over and over again, row by row from the top down. The image is also stored in row-major order in memory. The image contains n pixels, while the internal memory can hold M pixels and pixels are moved into and out of cache in blocks of size B , where $M \geq 3B$.

- (i) Construct an example—that is, pick values for B , M and n —such that the LRU caching policy results in M/B times more cache misses than an optimal caching policy for this application, excluding the first M/B cache misses of both strategies. (In other words, construct an example where the performance ratio of LRU versus optimal caching is M/B .)
- (ii) If we make the image only half the size—that is, n becomes half as large as in your example, but M and B stay the same—then what is the performance ratio of LRU versus optimal caching?
- (iii) If we make the image double the size, then what is the performance ratio of LRU versus optimal caching?

IO.II-2 (2 points) Consider an algorithm that needs at most cn/\sqrt{MB} I/Os if run with optimal caching, for some constant c . Prove that the algorithm needs at most $c'n/\sqrt{MB}$ I/Os when run with LRU caching, for a suitable constant c' .

Hint: Compare the performance of both versions to running the algorithm with optimal caching on a machine that has only half as much memory.

IO.II-3 (2+1 points) Let $A[0 \dots n-1]$ be an array of n distinct numbers. The *rank* of an element in $A[i]$ is defined as follows:

$$\text{rank}(A[i]) := (\text{number of elements in } A \text{ that are smaller than } A[i]) + 1.$$

Define the *displacement* of $A[i]$ to be $|i - \text{rank}(A[i]) + 1|$. Thus the displacement of $A[i]$ is equal to the distance between its current position in the array (namely i) and its position when A is sorted.

- (i) Suppose we know that A is already “almost” sorted, in the sense that the displacement of any element in A is less than $M - B$. Give an algorithm that sorts A using $O(n/B)$ I/Os, and argue that it indeed performs only that many I/Os.
- (ii) The $O(n/B)$ bound on the number of I/Os is smaller than the $\Omega((n/B) \log_{M/B}(n/B))$ lower bound from Theorem 6.2 from the Course Notes. Apparently the lower bound does not hold when the displacement of any element in A is less than $M - B$. Explain why the proof of Theorem 6.2 does not work in this case.
NB: You can keep your answer short—a few lines is sufficient—but you should point to a specific place in the proof where it no longer works.

IO.II-4 (0.5+1.5 points) Let $X[0..n-1]$ and $Y[0..n-1]$ be two arrays, each storing a set of n numbers. Let $Z[0..n-1]$ be another array, in which each entry $Z[i]$ has three fields: $Z[i].x$, $Z[i].y$ and $Z[i].sum$. The fields $Z[i].x$ and $Z[i].y$ contain integers in the range $\{0, \dots, n-1\}$; the fields $Z[i].sum$ are initially empty. We wish to store in each field $Z[i].sum$ the value $X[Z[i].x] + Y[Z[i].y]$. A simple algorithm for this is as follows.

```

ComputeSums( $X, Y, Z$ )
1: for  $i \leftarrow 0$  to  $n-1$  do
2:    $Z[i].sum \leftarrow X[Z[i].x] + Y[Z[i].y]$ 
3: end for

```

- (i) Analyze the number of I/Os performed by *ComputeSums*.
- (ii) Give an algorithm to compute the values $Z[i].sum$ that performs only $O(\text{SORT}(n))$ I/Os.

Exercise Set I/O-Efficient III

IO.III-1 ($1\frac{1}{2} + 1 + 1\frac{1}{2}$ points) Consider a complete binary search tree \mathcal{T} with $n = 2^k - 1$ nodes—that is, a binary search tree with $k = \lceil \log n \rceil$ levels that are all completely filled—which is stored in external memory. In this exercise we investigate the effect of different blocking strategies for the nodes of \mathcal{T} , where we assume that the internal memory can hold at least two blocks.

- (i) Suppose blocks are formed according to an in-order traversal of \mathcal{T} (which is the same as the sorted order of the values of the nodes). Analyze the minimum and maximum number of I/Os needed to traverse any root-to-leaf path in \mathcal{T} .
- (ii) Describe an alternative way to form blocks, which guarantees that any root-to-leaf path can be traversed in $O(\log_B n)$. What is the relation of your blocking strategy to B-trees?
- (iii) Prove that for *any* blocking strategy—that is, any strategy to group the nodes into blocks with at most B nodes each—there is a root-to-leaf path that visits $\Omega(\log_B n)$ nodes.

IO.III-2 (3 points) Let S be an initially empty set of numbers. Suppose we have a sequence of n operations op_0, \dots, op_{n-1} on S . Each operation op_i is of the form $(type_i, x_i)$, where $type_i \in \{\text{Insert}, \text{Delete}, \text{Search}\}$ and x_i is a number. You may assume that when a number x is inserted it is not present in S , and when a number x is deleted it is present. (After a number has been deleted, it could be re-inserted again.) The goal is to report for each of the *Search*-operations whether the number x_i being searched for is present in S at the time of the search. With a buffer tree these operations can be performed in $O(\text{SORT}(n))$ I/Os in total. Show that the problem can be solved more directly (using sorting) in $O(\text{SORT}(n))$ I/Os, without using a buffer tree.

IO.III-3 (3 points) Let $\mathcal{G} = (V, E)$ be an undirected graph stored in the form of an adjacency list in external memory. A *minimal vertex cover* for \mathcal{G} is a vertex cover C such that no vertex can be deleted from C without losing the cover property. (In other words, there is no $v \in C$ such that $C \setminus \{v\}$ is also a valid vertex cover). Prove that it is possible to compute a minimal vertex cover for \mathcal{G} using only $O(\text{SORT}(|V| + |E|))$ I/Os.